

Simulating seismic wave propagation with sw4

CIG tutorial

N. Anders Petersson

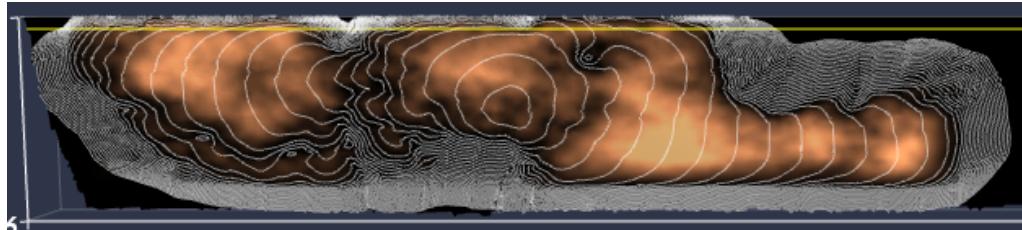
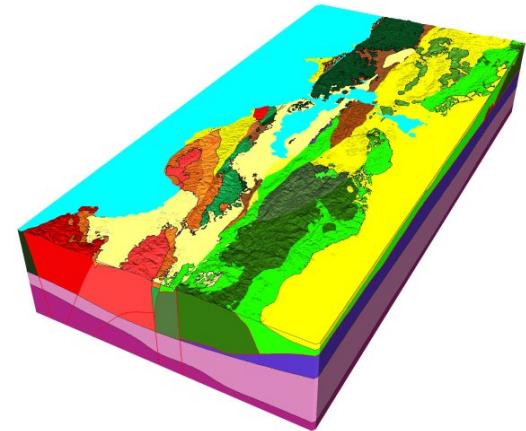
June, 2016



SW4 simulates regional earthquake ground motions on parallel computers

- Mathematical model of the Earth:
 - Navier's eqn of 3-D linear isotropic solid dynamics
 - 4th order summation-by-parts finite differences
 - 4th order time explicit stepping
- Geophysical model
 - Isotropic heterogeneous elastic & visco-elastic material
 - Topography
- Kinematic source rupture model

$$\rho \frac{\partial^2 \mathbf{u}}{\partial t^2} = \nabla \cdot \mathcal{T} + \mathbf{f}$$

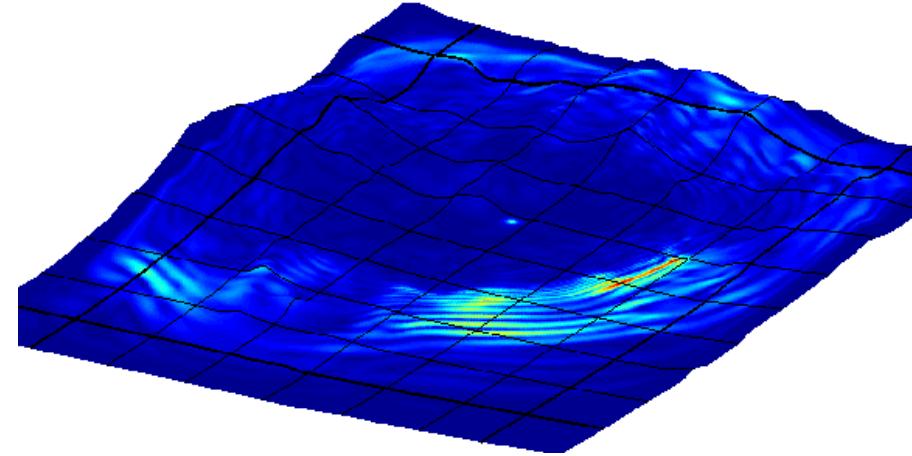


- MPI-based distributed memory programming model
 - C++ with Fortran numerical kernels
- SW4 is open source, GNU GPL-v2



The SW4 and WPP codes solve the time dependent 3D seismic wave equation

- WPP is 2nd order accurate and SW4 is 4th order
- Summation by parts (SBP) finite difference methods
 - Fundamentally different from FDTD
 - Energy stability
 - Higher order accuracy
- Curvilinear coordinates
 - Realistic topography
 - Absorbing far-field layers
- Visco-elastic modeling
- Point forces and moment tensor sources
 - Grid independent location
- Anisotropic materials
- Local mesh refinement



Summation by parts discretization mimic integration by parts for finite differences

- Elastic wave equation:

$$\rho \frac{\partial^2 \mathbf{u}}{\partial t^2} = \mathbf{L}\mathbf{u} + \mathbf{f}(\mathbf{x}, t)$$

$$\mathbf{L}\mathbf{u} := \nabla \cdot \mathcal{T}$$

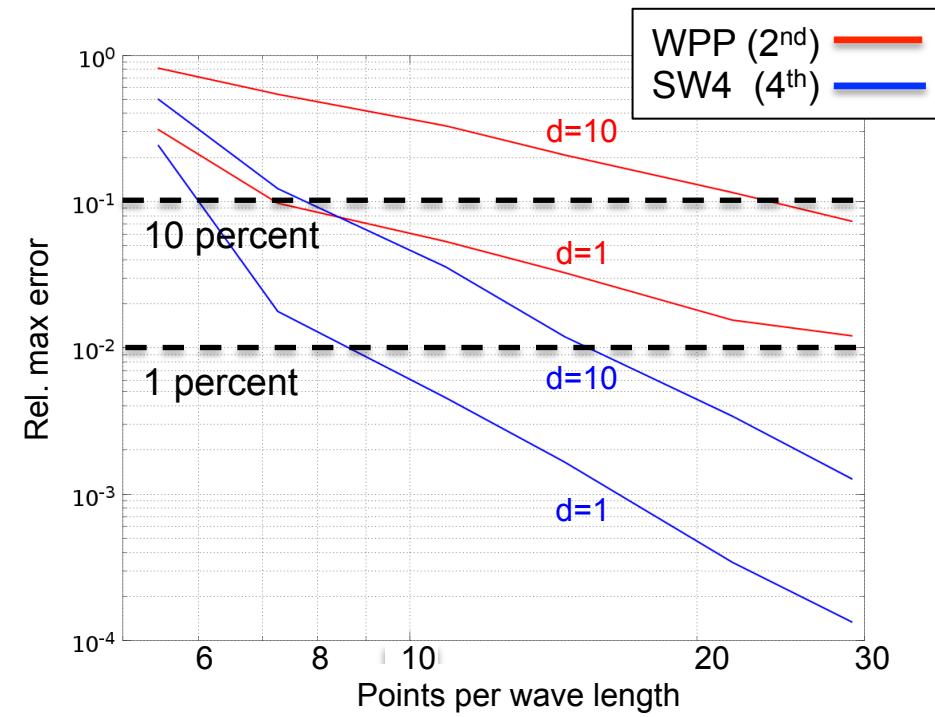
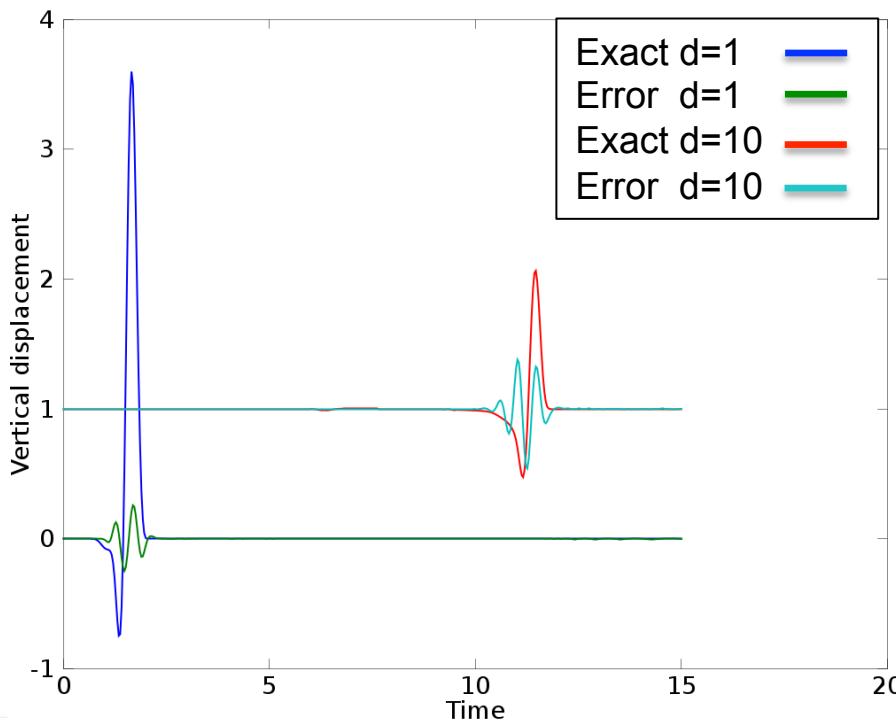
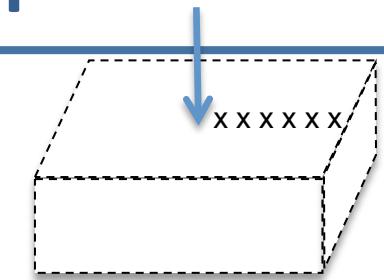
- Discretized in space:

$$\rho \frac{d^2 \mathbf{u}}{dt^2} = \mathbf{L}_h \mathbf{u} + \mathbf{f}(\mathbf{x}, t)$$

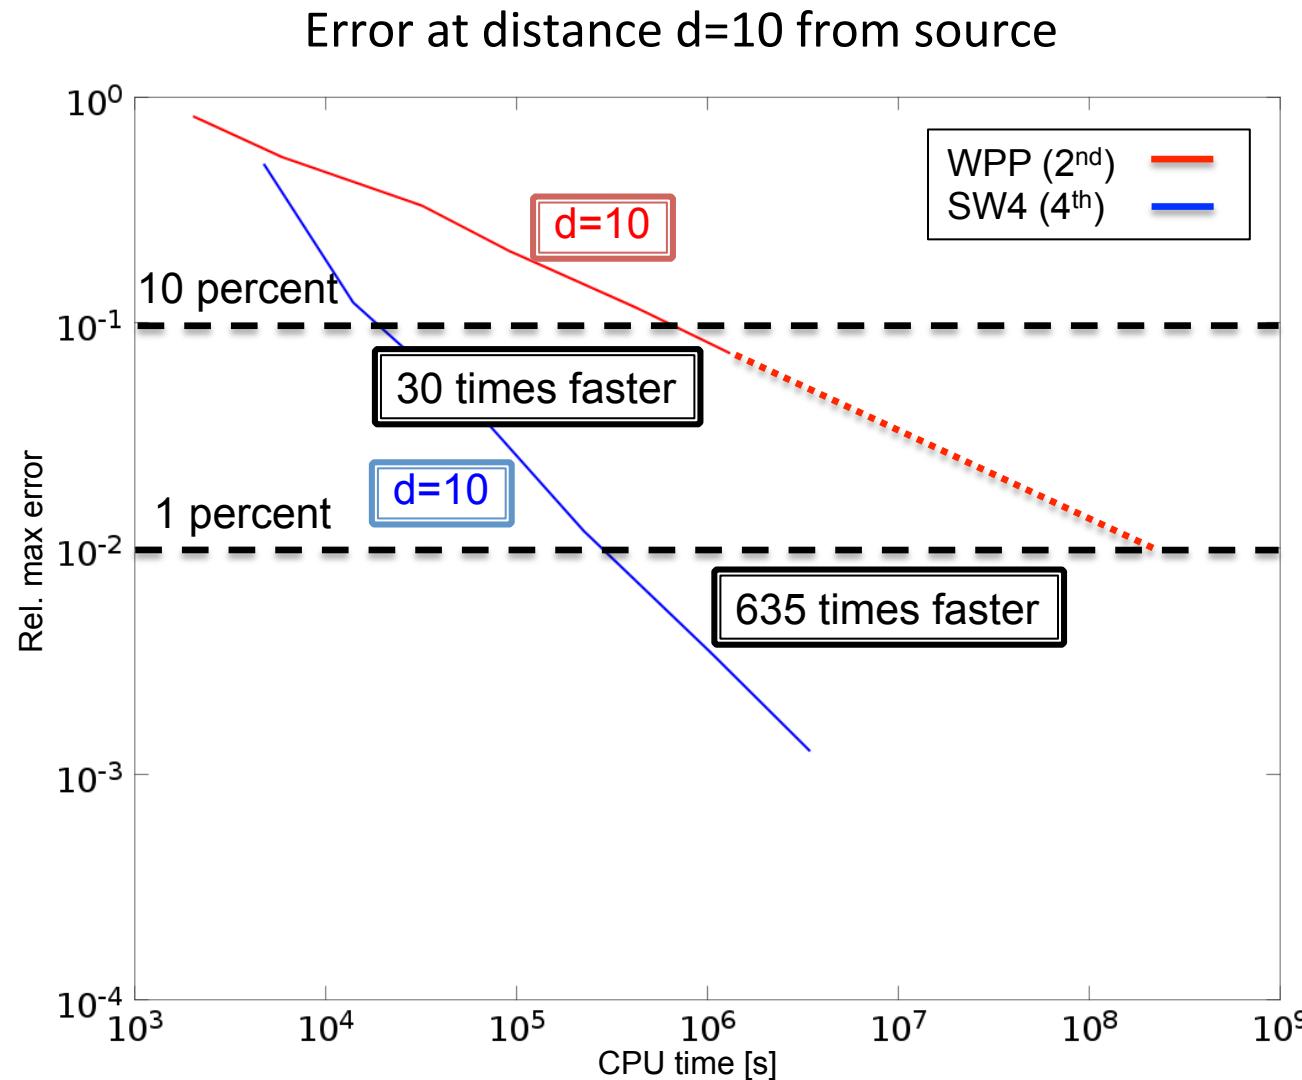
- Symmetry and negative definite properties of ‘ \mathbf{L} ’ preserved by summation by parts discretization ‘ \mathbf{L}_h ’
- Numerical scheme is energy stable

A 4th order method is more efficient since it needs fewer grid points per wavelength

- Lamb's problem
- Solution at distance $d=1$ and $d=10$ [Mooney (1974)]
- Error: P^{-2} versus P^{-4} (P = grid points per wave length)
- Improved efficiency with 4th order accuracy

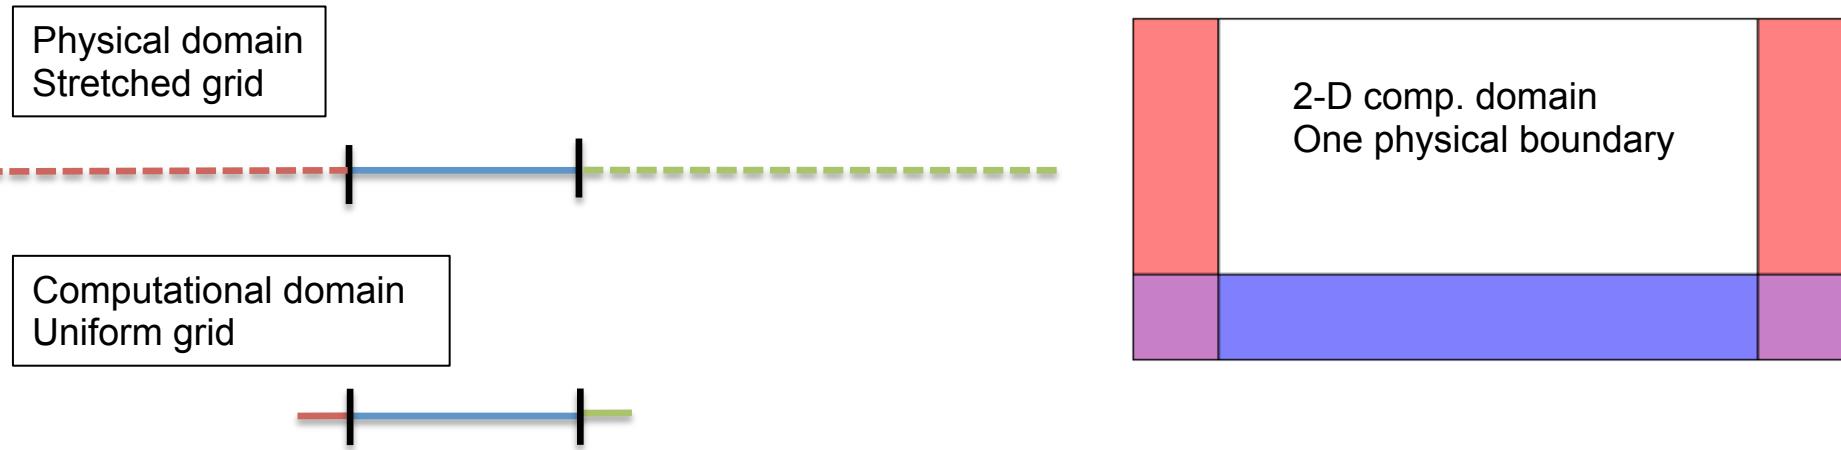


The 4th order method is much faster at larger distances and gives smaller errors

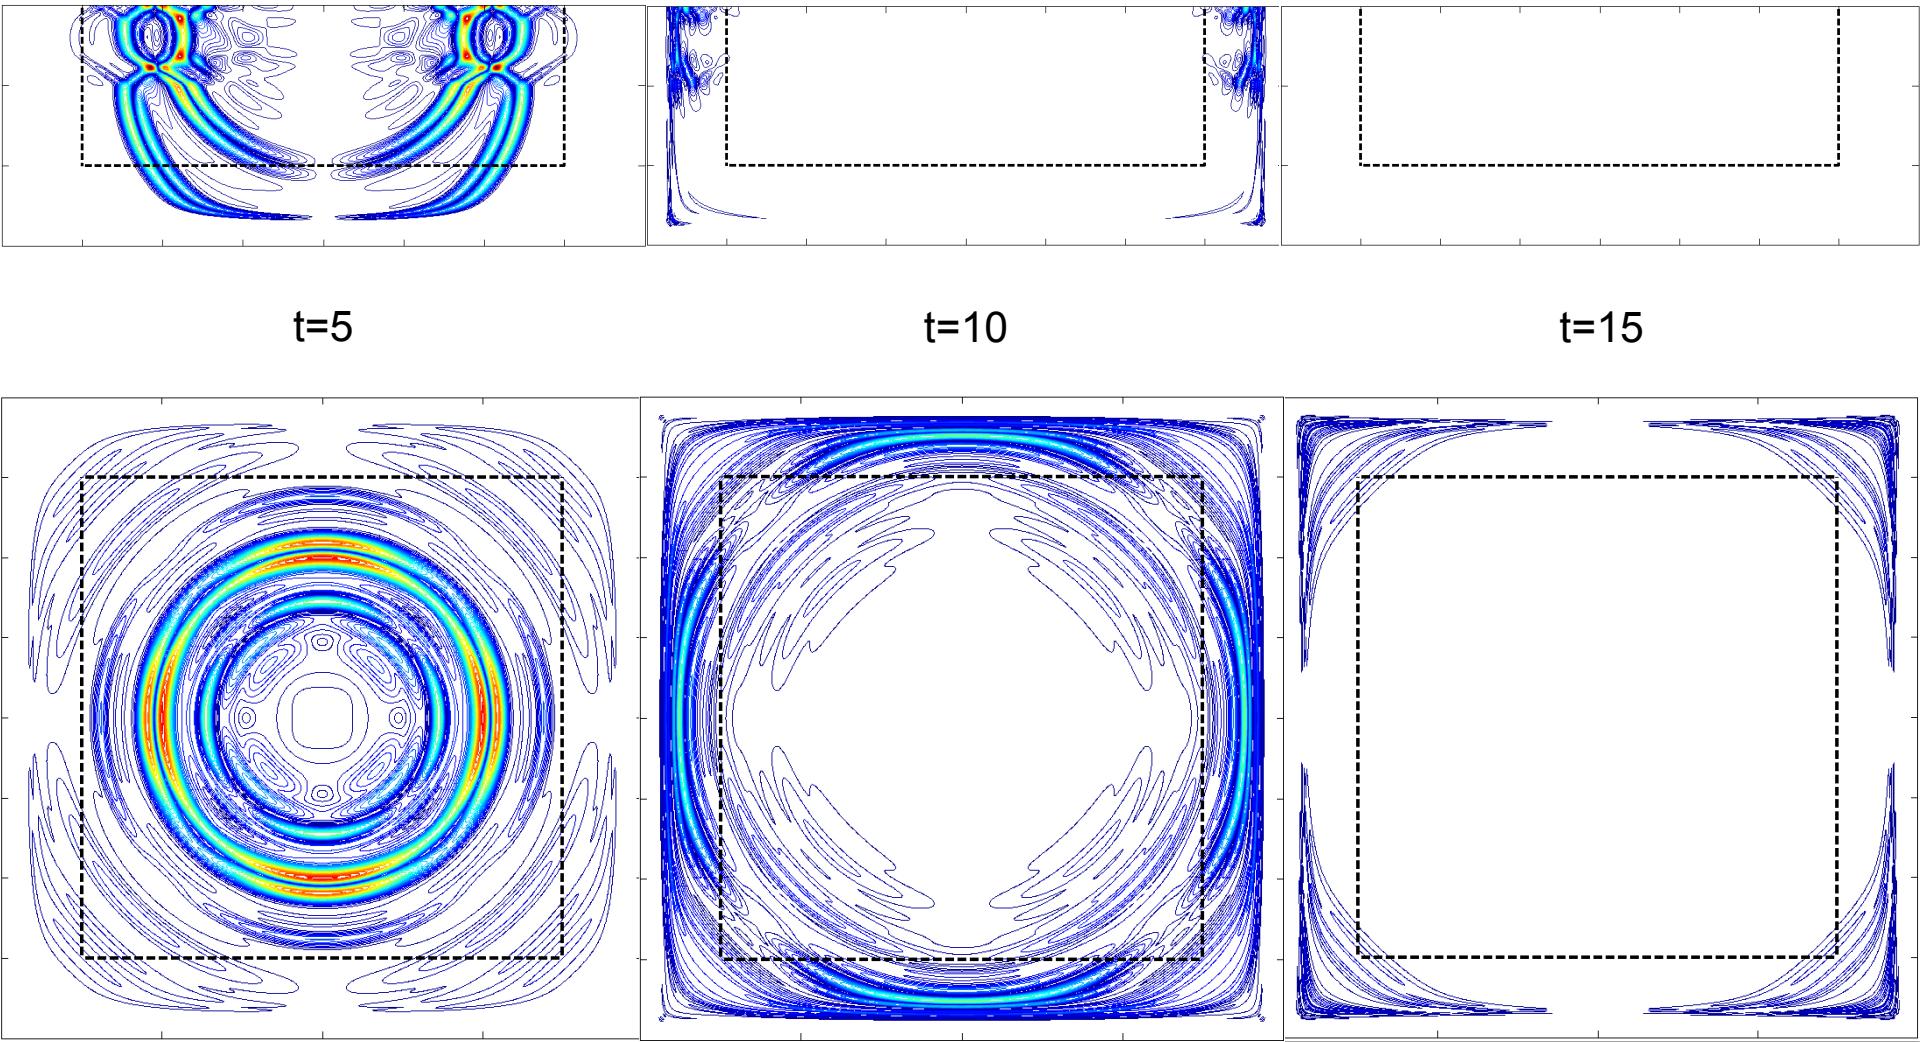


SW4 uses super-grid layers to suppress artificial reflections from far-field boundaries

- Perfectly matched layers (PML) can be unstable
- Super-grid shrinks very large physical domain
 - Slow down and compress waves in layers
 - Damp unresolved modes with high order dissipation
- Always stable
 - Accurate if layer is sufficiently wide

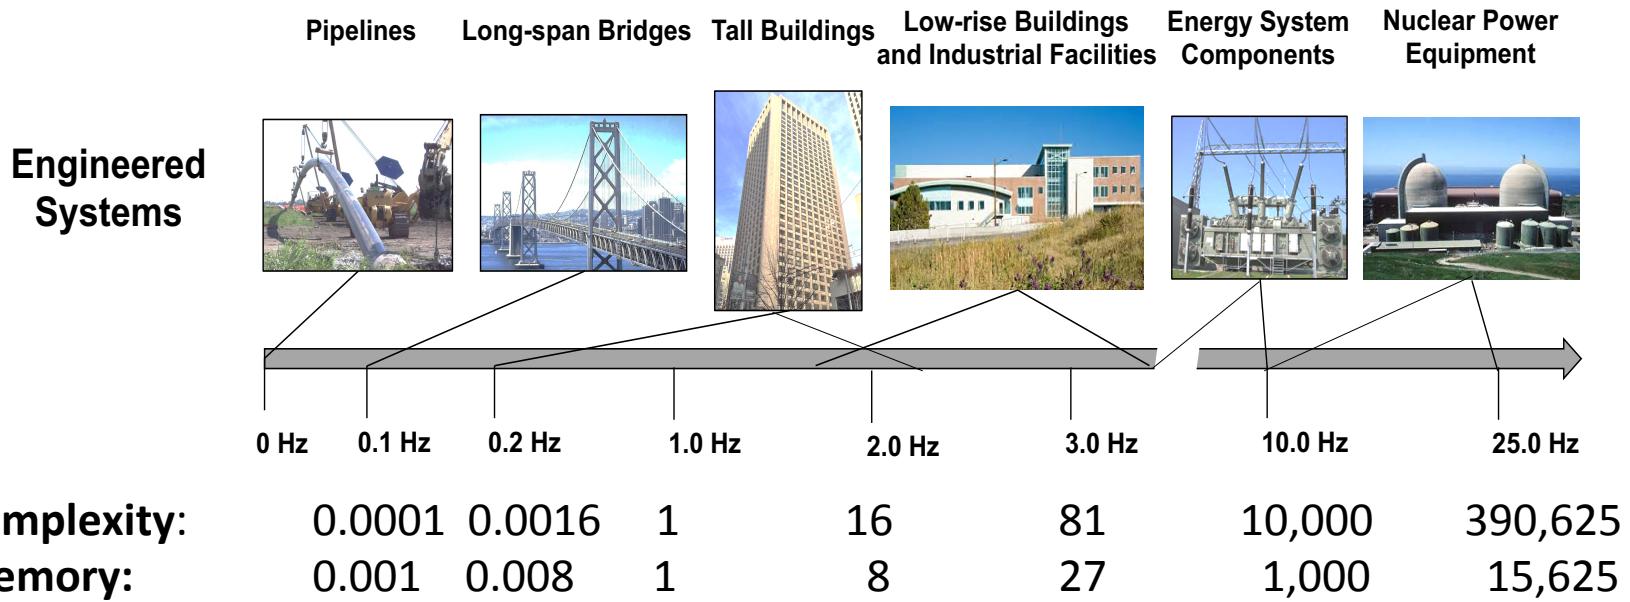


Super-grid test: 3-D elastic wave equation with layered material and a free surface



The computational requirements scale as (frequency)⁴

- Smallest wave length $\lambda \approx V_s/f_{\max}$
- Grid size: $h = \lambda/P$; Points per wavelength: P
- Double frequency $\rightarrow 8 \times$ grid points, $16 \times$ CPU



Hands-on sw4 tutorial

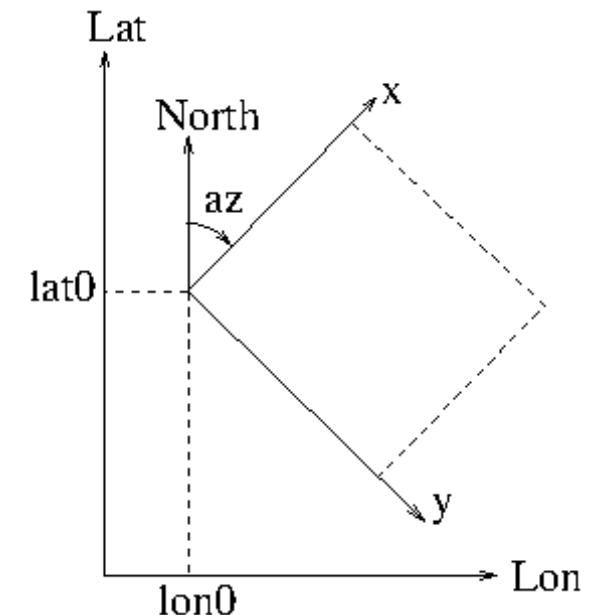
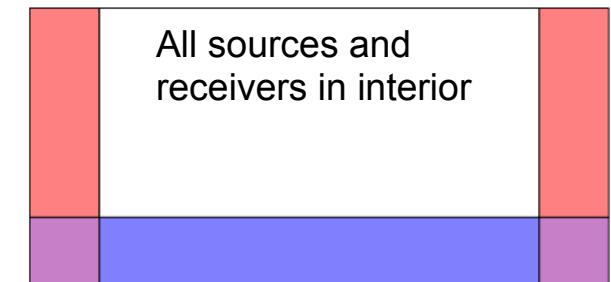
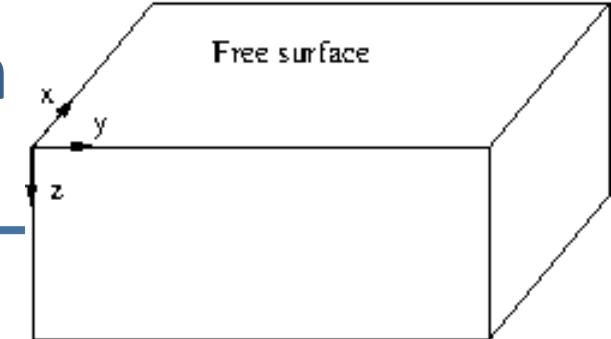
- How to run sw4
 - Smaller & larger jobs
- How to set up the simulation
 - Sources
 - How do I pick the grid size?
 - Topography
 - Material model, visco-elastic modeling
 - Output options
- Practical suggestions
 - Start coarse, check model, estimate resources, refine
 - Use the parallel file system
 - Initial analysis before moving large files
- Lots of info in the user's guide
 - VM: ~/Documents/SW4-UsersGuide.pdf

Building and running sw4

- sw4 can be built on Linux/OSX/Unix machines
 - ~/Documents/SW4-Installation.pdf
- Command file specifies the simulation
 - mpirun -np 2 sw4 **my-input.in**
- Syntax for running MPI-jobs varies:
 - Macs with openmpi: openmpirun -np 4 sw4 input.in
 - VM-ware mpirun: mpirun -np 8 sw4 input.in
 - Livermore computing (LC): srun -n4096 sw4 my-big-job.in
- Running on supercomputers (details may vary)
 - Smaller, shorter jobs on interactive debug partition. Start job from command prompt. No CPU-bank needed
 - Larger jobs use a batch system. Job described by a script. Usually needs a CPU-bank
 - For faster I/O, use a parallel file system if available

sw4 uses SI-units (MKS), angles in degrees

- Box-shaped computational domain
 - Right-handed system, z positive downwards
- Top surface optionally follows topography
- Free surface boundary condition on top surface
- Super-grid absorbing layers on all other sides
 - SG layers included in computational domain
 - Default thickness: 30 grid points
 - Smallest Cartesian grid: 60x60x30
- Locations in Cartesian or geographic coordinates
- Geographic coordinates are mapped:
 - Default: spheroidal projection
 - Proj.4 library for better accuracy
- Rotate grid with azimuth angle (az)
 - az=0: x-axis = North, y-axis = East



1st exercise: solve Lamb's problem

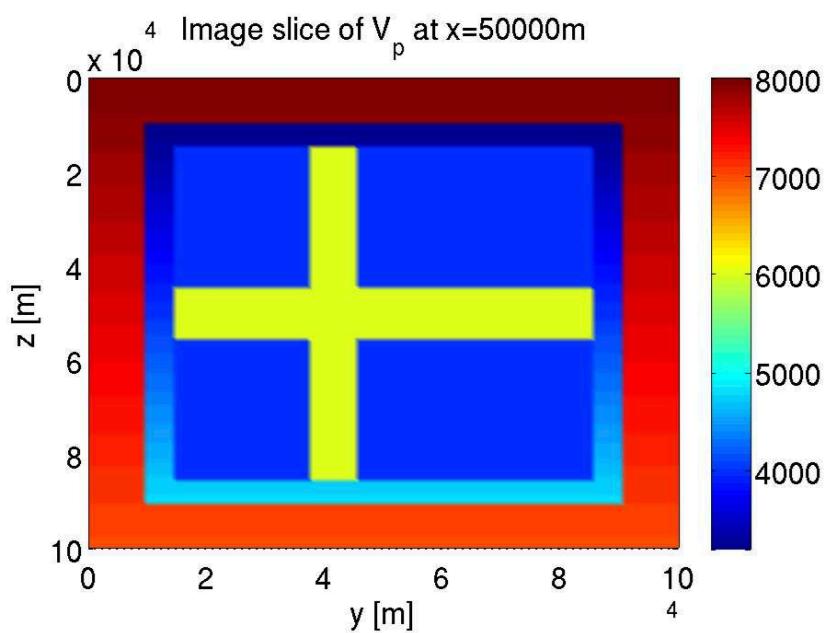
- Start VM virtualBox & import the CIG16-TutorialVM appliance
- Log on as sw4_user (No password)
- cd Desktop
- Simulation described in input file: seismic1.in
 - fileio path=seis-out
 - grid nx=151 x=10e3 y=10e3 z=5e3
 - time t=5.0
 - # supergrid thickness for h = 10000/150 = 1000/15, 30*h = 2000, 60*h = 4000
 - supergrid gp=60
 - block vp=1.7320508076e+03 vs=1000 rho=1500
 - source type=C6SmoothBump x=5e3 y=5e3 z=0 fz=1e13 freq=1 t0=0
 - # Time history of solution on the surface
 - rec x=5e3 y=6e3 z=0 file=v1s sacformat=0 usgsformat=1
- mpirun -np 1 sw4 seismic1.in
- (1 min, 51 sec runtime on my Mac laptop)

We use octave to view the results

- octave is open source software (similar to Matlab)
- Open octave from terminal window (in VM):
 - >octave
- Expand octave search path (put command in `~/.octaverc` file)
- `octave:1> addpath("/home/sw4_user/Desktop/tools")`
- `octave:2>cd Desktop/seis-out`
- `octave:3>ls`
- `octave:4>plotusgs("v1s.txt")`

The material model (ρ , C_p , C_s) must be defined for all points in the grid

- One or several material commands can be combined
 - block, rfile, efile, pfile, ifile, ...
- The order matters (later commands take precedence)
- Visco-elastic modeling triggered by attenuation command:
Also need to specify Q_p and Q_s

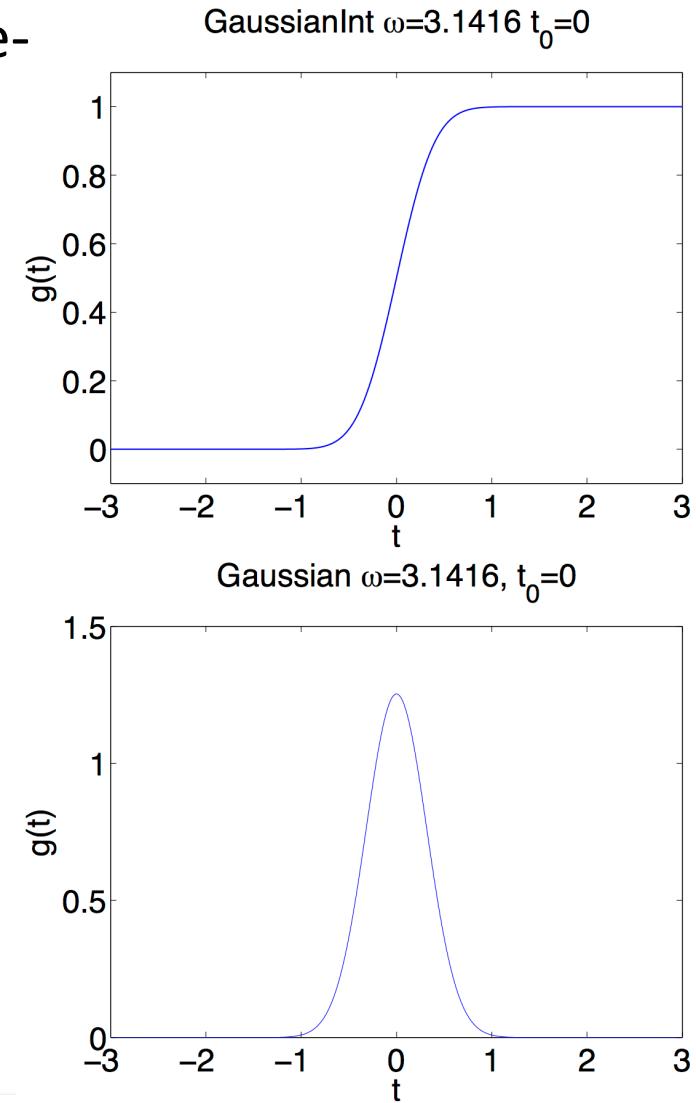


sw4 provides several options for specifying sources and time-functions

- Any number of point forces or moment tensor sources
 - Forces defined by 3-component vector
 - Moment tensor (symmetric) defined by 6 components M_{ij} , or
 - Seismic moment (M_0) and strike, dip, rake angles
- Many pre-defined time functions (Gaussian, Ricker, Liu, ...)
 - Frequency parameter and start/center time
- Dirac delta-distribution: trigger all frequencies on mesh
 - Discrete Green's functions. Motion must be filtered for accuracy
- User defined discrete time function (interpolated by spline)
- Complex ruptures can be specified by using many point sources or an SRF file

There are several ways of computing displacements, velocities, or acceleration

- Displacements correspond to source time-functions that tends to a constant (e.g. GaussianInt)
 - Displacements obtained directly
 - Velocities by differentiating once
 - Acceleration by differentiating twice
- Velocities correspond to time-functions that tend to zero but have non-zero integral (e.g. Gaussian)
 - Velocities obtained directly
 - Displacements by integrating once
 - Acceleration by differentiating once



The highest frequency and the lowest wave speed determine the grid size

- Shortest wave length $L_{\min} = \min C_s / f_{\max}$
- Grid points per shortest wave length $P_{\min} = L_{\min} / h$
- Good accuracy when $P_{\min} > 6$ to 10 (depending on distance)
- 3 approaches:
 - Tune freq in source time function to not trigger unresolved waves
 - Calculate $f_{\max} = \min C_s / (P_{\min} h)$, filter out unresolved frequencies afterwards (sac)
 - Use prefilter command to pre-process all sources (not discussed today)
- Relation between f_{\max} and frequency parameter (freq) is different for each time function (see UG 4.4)
- octave scripts for plotting time function and Fourier transform: `fcnplot`, `ftfcnplot`

2nd exercise: solve a layer over half-space problem. Input file: loh1-slow.in

- The order of the block commands is important

```
block vp=6000 vs=3464 rho=2700
```

```
block vp=4000 vs=2000 rho=2600 z2=1000
```

- Gaussian time function gives the velocities directly

```
source x=15000 y=15000 z=2000 mxy=1e18 t0=1.2 freq=5.0 type=Gaussian
```

- Time histories recorded @ the closest grid point

```
rec x=15600 y=15800 z=0 file=sta01 usgsformat=1
```

```
rec x=21000 y=23000 z=0 file=sta10 usgsformat=1
```

- Material images don't change, recorded initially

```
image mode=s y=15e3 file=ver cycle=0
```

- Solution images can be recorded at fixed time intervals

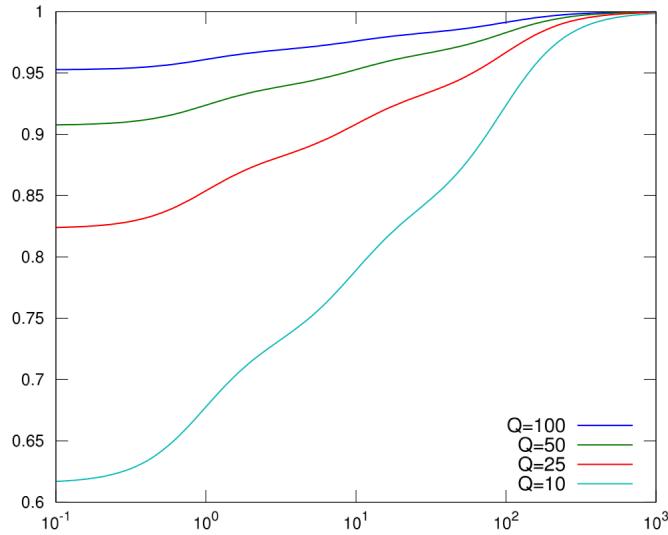
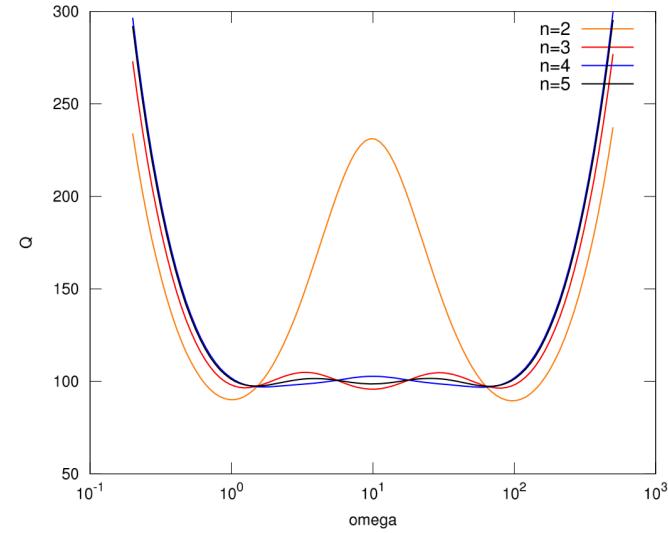
```
image mode=mag z=0 file=surf timeInterval=0.5
```

We can use octave to view image files

- octave>cd Desktop/loh1-slow
- octave>ls
- octave>help plotimage
- octave>plotimage("surf.cycle=170.z=0.mag.sw4img")
- octave>help imageinfo
- octave>imageinfo("surf.cycle=170.z=0.mag.sw4img",10,1)

Standard linear solid (SLS) mechanisms are used to model anelastic attenuation

- The quality factors (Q_P and Q_S) are approximately constant in frequency band $[0.01 f_{\max}, f_{\max}]$
 - User must specify upper frequency: f_{\max}
 - More mechanisms: less variation in Q, but more calculations. Default nmech=3
- The visco-elastic material is dispersive: phase velocity depends on frequency
 - User must specify phasel freq parameter: frequency where C_P and C_S are accurate.



3rd exercise: solve an anelastic layer over half-space problem. Input file: loh3-slow.in

- Only differences from loh1-slow.in

different output directory

fileio path=loh3-slow

must specify Qp and Qs

block vp=6000 vs=3464 rho=2700 Qs=69.3 Qp=155.9

block vp=4000 vs=2000 rho=2600 z2=1000 Qs=40 Qp=120

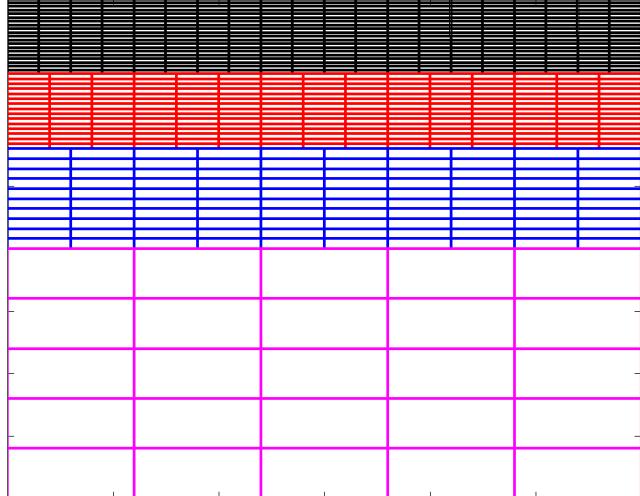
enable attenuation modeling

attenuation phasefreq=2.5 nmech=3 maxfreq=15

- About 3x CPU time and 2.7x memory

Large complex material models can be described in an rfile model

- Variable resolution data structure
- 7.6 Gbyte file for regional SF bay area
- Vulcan (5 PFlops): 1 Gbyte memory per core
- Parallel I/O routines on parallel file systems

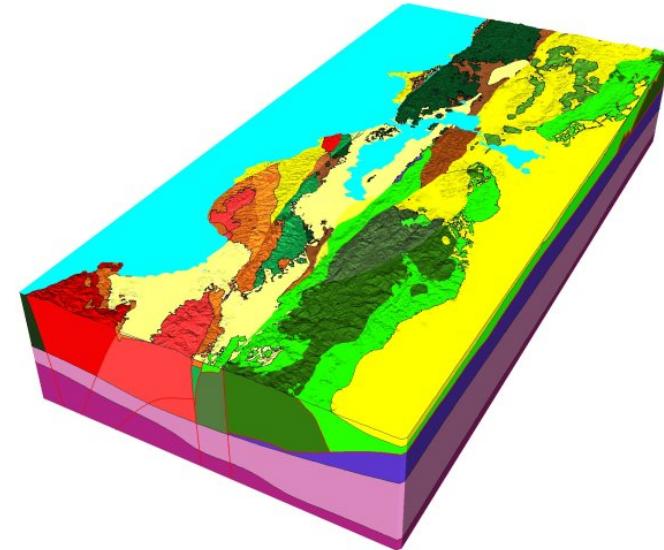


2897 x 1401 x 74 (6 Gb)

1449 x 701 x 57 (1.2 Gb)

725 x 351 x 33 (0.17 Gb)

363 x 176 x 194 (0.25 Gb)



Realistic topography is included in an rfile

- SW4 constructs a curvilinear grid on startup
- $z=0$ is mean sea level
- User must pick bottom z-coordinate for curvilinear grid: z_{max}
- Rule of thumb: $z_{max} = \text{twice the variation below the lowest elevation}$
 - $1500 < \text{elevation} < 2500 \text{ m (positive above sea level)}$
 - $z_{max} \geq - (1500 - 2 * (2500 - 1500)) = +500 \text{ (positive below sea level)}$



4th exercise: Simulate earthquake near Berkeley.

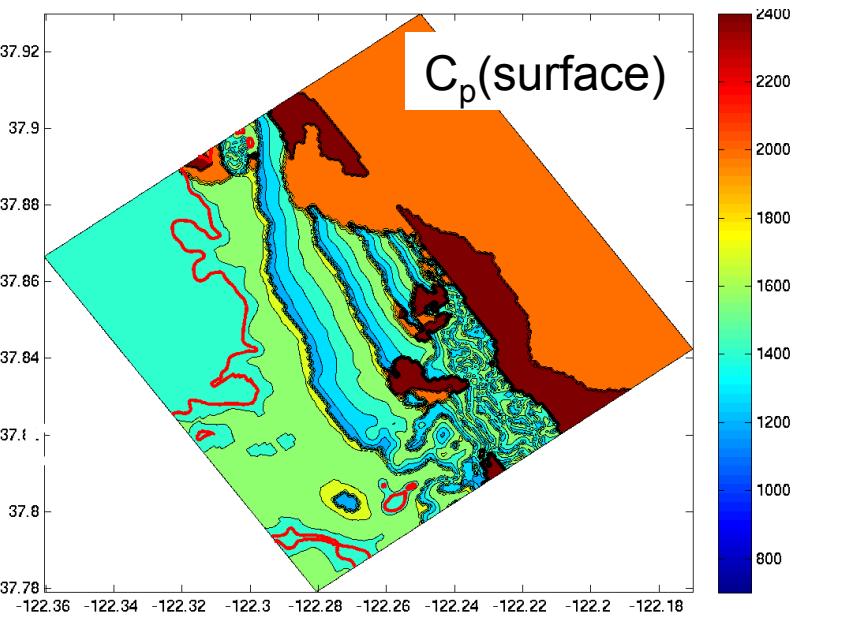
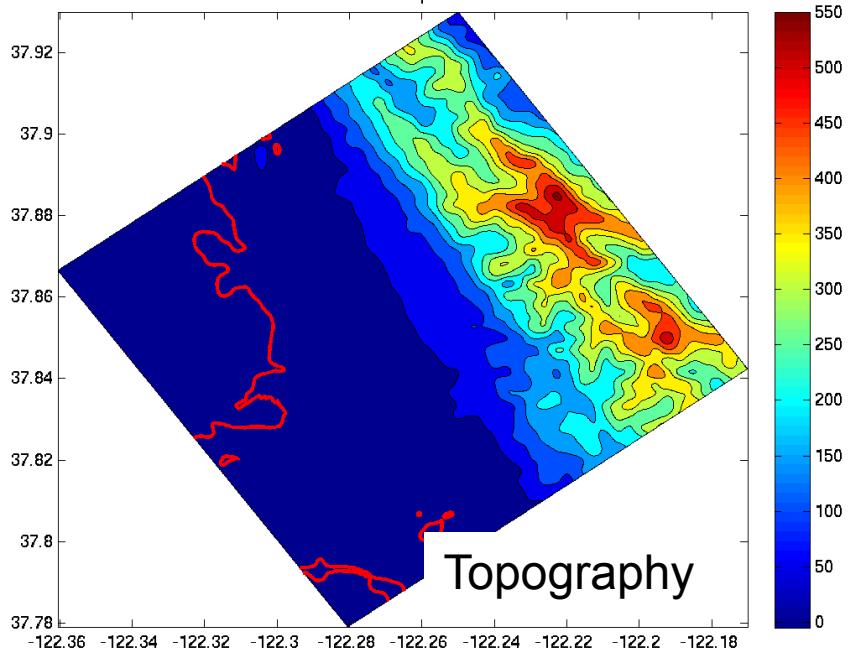
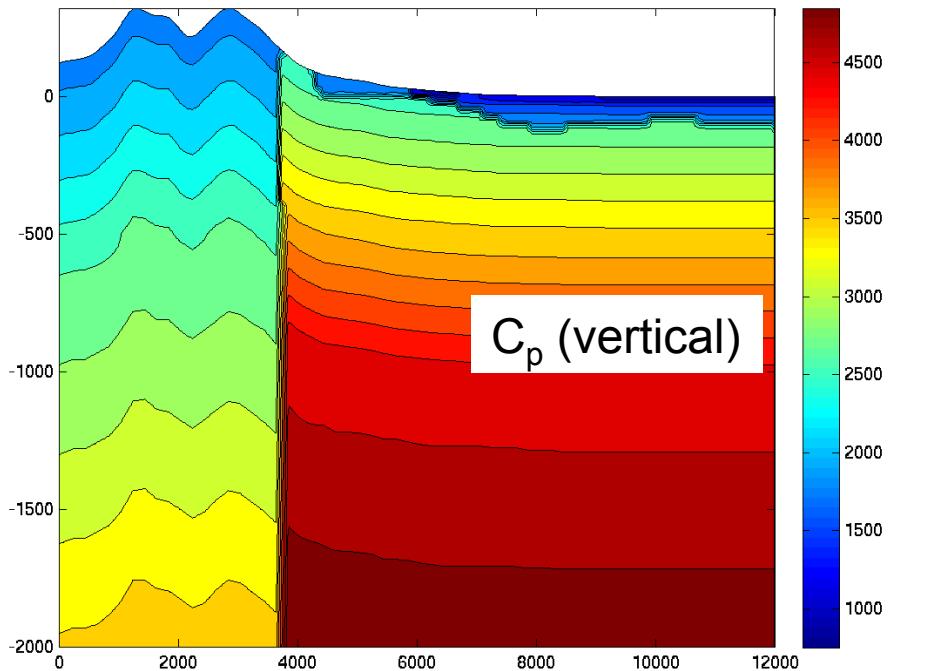
Input file: berkeley-small.in

- Proj.4 commands in grid command (az must agree with rfile header)

```
grid x=12e3 y=12e3 z=5e3 nx=151 lat=37.93 lon=-122.25 az=143.638  
proj=tmerc datum=NAD83 lon_p=-123.0 lat_p=35.0 scale=0.9996
```

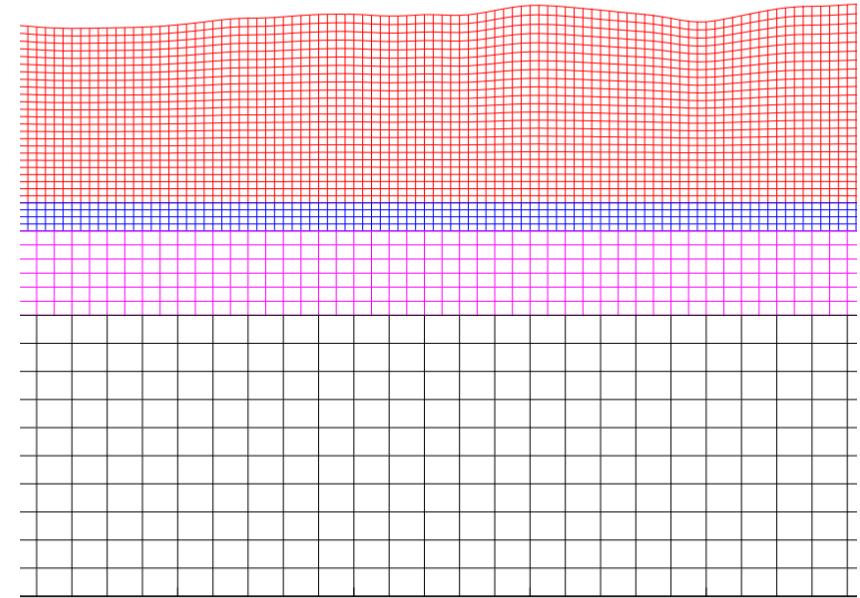
- Topography command
topography input=rfile zmax=2e3 order=3 file=berkeley-r.rfile
- Material model
rfile filename=berkeley-r.rfile
- Runtime 9 min, 50 sec

The case berkeley-small.in uses a small rfile model



Mesh refinement can be used to save computational resources

- Keep number of grid points per wave length \sim constant
- Saves memory
- Allows larger time step
- Currently only elastic materials



5th exercise: Try mesh refinement. Input file meshref.in

- The grid command specifies the coarsest grid size
`grid h=350 x=30000 y=30000 z=17000`
- User specifies z-level with the refinement command
`refinement zmax=2.5e3`
- Each grid patch must have at least 12 grid points in the vertical direction.
- Runtime ~5 min
- Same problem w/o mesh refinement in nomeshref.in

General guidelines

- Check setup on a coarse grid
- Check a few images of the material model
- Check receivers on a map: gmt
- Extrapolate computational resources from coarse run
- Doubling frequency = Halving the grid size = doubling # grid points / dimension
- Doubling # grid points/dimension: 8x grid points, 2x time steps
 - On same number of cores: 8x memory, 16x CPU time
 - On 8x more cores: same memory, 2x CPU time (weak scaling)
 - On 16x cores: $\frac{1}{2}$ x memory, about the same CPU time
- SW4 source distribution includes examples and octave scripts

What if it doesn't work?

- Check command syntax and rules in SW4 user's guide
- Update SW4 source code from github
- Ask questions on CIG-SEISMO mailing list

cig-seismo@geodynamics.org

- Discuss problems in the github issue tracker

<https://github.com/geodynamics/sw4/issues>

- Help provided as time permits



**Lawrence Livermore
National Laboratory**